

LAB3: BINARY CLASSIFICATION OF HEART DISEASE OF PATIENTS USING DEEP NEURAL NETWORK

In [1]: `import pandas as pd`

In [2]: `df = pd.read_csv("heart_data.csv")`
`df.head(10)`

Out[2]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
5	57	1	0	140	192	0	1	148	0	0.4	1	0	1	1
6	56	0	1	140	294	0	0	153	0	1.3	1	0	2	1
7	44	1	1	120	263	0	1	173	0	0.0	2	0	3	1
8	52	1	2	172	199	1	1	162	0	0.5	2	0	3	1
9	57	1	2	150	168	0	1	174	0	1.6	2	0	2	1

In [3]: `df.shape`

Out[3]: (303, 14)

In [4]: `df.size`

Out[4]: 4242

In [5]: `df.columns`

Out[5]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'], dtype='object')

2. Split the dataset for training and testing (test size = 20%)

In [6]: `X = df.drop('target',axis=1)`
`y = df['target']`

In [8]: X

Out[8]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2

303 rows × 13 columns

In [13]: y

Out[13]:

0	1
1	1
2	1
3	1
4	1
...	..
298	0
299	0
300	0
301	0
302	0

Name: target, Length: 303, dtype: int64

In [14]: `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s`

In [15]: X_train.shape

Out[15]: (242, 13)

In [12]: X_test.shape

Out[12]: (61, 13)

3. Create a neural network based on the following requirements

```
In [16]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
In [17]: model = Sequential()
model.add(Dense(8, input_dim=13, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
```

4. Compile your model with learning rate = 0.001, optimizer as 'RMSprop', Mean square error loss and metrics as 'accuracy'.

```
In [18]: from tensorflow import keras
```

```
In [19]: optimizer = keras.optimizers.RMSprop(learning_rate=0.001)
```

```
In [20]: model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 3s 9ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 2/10
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 3/10
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 4/10
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 5/10
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 6/10
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 7/10
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 8/10
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 9/10
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 10/10
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy: 0.4504
```

```
Out[20]: <keras.callbacks.History at 0x1d7869cf0d0>
```

In [21]: `model.evaluate(X_test, y_test)`

2/2 [=====] - 0s 3ms/step - loss: 0.5246 - accuracy: 0.4754

Out[21]: [0.5245901346206665, 0.4754098355770111]

5. Print the summary of the model: `model.summary()`

In [22]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8)	112
dense_1 (Dense)	(None, 1)	9
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		

6. Train the model for 200 epochs and batch size as 10

In [23]: `model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])`
`model.fit(X_train, y_train, epochs=200, batch_size=10, verbose=1)`

25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 13/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 14/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 15/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 16/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 17/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 18/200
25/25 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy: 0.4504
Epoch 19/200

In [24]: `model.evaluate(X_test, y_test)`

2/2 [=====] - 0s 2ms/step - loss: 0.5246 - accuracy: 0.4754

Out[24]: [0.5245901346206665, 0.4754098355770111]

7. Save the trained model in a variable, such as, history. Also, you can split your training data for validation such as 20% of training data.

In [26]: `history = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=32)`

Epoch 1/100
20/20 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 2/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 3/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 4/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 5/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 6/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286

8. Evaluate the trained model to predict the probability values for the test data set (ie., xtest and ytest)

In [27]: `model.evaluate(X_test, y_test)`

2/2 [=====] - 0s 3ms/step - loss: 0.5246 - accuracy: 0.4754

Out[27]: [0.5245901346206665, 0.4754098355770111]


9. Print the model accuracy and model loss as below (Use can use the 'history' object we have saved). Sample code is given below.


In [28]: `history.history.keys()`

Out[28]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [31]: `import matplotlib.pyplot as plt`

```
In [32]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
```





10. Do further experiments

```
In [33]: model1 = Sequential()  
model1.add(Dense(16, input_dim=13, activation='relu'))  
model1.add(Dense(8, activation='relu'))  
model1.add(Dense(1, activation='sigmoid'))
```

```
In [34]: model1.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
model1.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 1ms/step - loss: 0.4397 - accuracy:  
0.5207  
Epoch 2/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3482 - accuracy:  
0.5868  
Epoch 3/10  
9/9 [=====] - 0s 2ms/step - loss: 0.3434 - accuracy:  
0.5992  
Epoch 4/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3535 - accuracy:  
0.5992  
Epoch 5/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3287 - accuracy:  
0.6157  
Epoch 6/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3256 - accuracy:  
0.6198  
Epoch 7/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3232 - accuracy:  
0.6157  
Epoch 8/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3080 - accuracy:  
0.6446  
Epoch 9/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2986 - accuracy:  
0.6612  
Epoch 10/10  
9/9 [=====] - 0s 1ms/step - loss: 0.3306 - accuracy:  
0.6198
```

```
Out[34]: <keras.callbacks.History at 0x1d78bda9d80>
```

```
In [35]: model1.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.2617 - accuracy:  
0.7049
```

```
Out[35]: [0.26170825958251953, 0.7049180269241333]
```



```
In [36]: history1 = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=32)
0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 5/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 6/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 7/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 8/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 9/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 10/100
7/7 [=====] - 0s 6ms/step - loss: 0.5440 - accuracy: 0.4560 - val_loss: 0.5714 - val_accuracy: 0.4286
Epoch 11/100
```

```
In [37]: model1.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 16)	224
dense_3 (Dense)	(None, 8)	136
dense_4 (Dense)	(None, 1)	9
Total params: 369		
Trainable params: 369		
Non-trainable params: 0		

```
In [38]: ls = history1.history
```

```
In [39]: new = pd.DataFrame.from_dict(ls)
new
```

Out[39]:

	loss	accuracy	val_loss	val_accuracy
0	0.544041	0.455959	0.571429	0.428571
1	0.544041	0.455959	0.571429	0.428571
2	0.544041	0.455959	0.571429	0.428571
3	0.544041	0.455959	0.571429	0.428571
4	0.544041	0.455959	0.571429	0.428571
...
95	0.544041	0.455959	0.571429	0.428571
96	0.544041	0.455959	0.571429	0.428571
97	0.544041	0.455959	0.571429	0.428571
98	0.544041	0.455959	0.571429	0.428571
99	0.544041	0.455959	0.571429	0.428571

100 rows × 4 columns

```
In [40]: model2 = Sequential()
model2.add(Dense(32, input_dim=13, activation='relu'))
model2.add(Dense(16, activation='relu'))
model2.add(Dense(8, activation='relu'))
model2.add(Dense(1, activation='sigmoid'))
```

```
In [41]: model2.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model2.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 1s 2ms/step - loss: 0.4083 - accuracy: 0.4711
Epoch 2/10
9/9 [=====] - 0s 2ms/step - loss: 0.3643 - accuracy: 0.4463
Epoch 3/10
9/9 [=====] - 0s 2ms/step - loss: 0.3247 - accuracy: 0.4421
Epoch 4/10
9/9 [=====] - 0s 2ms/step - loss: 0.3017 - accuracy: 0.5083
Epoch 5/10
9/9 [=====] - 0s 2ms/step - loss: 0.2687 - accuracy: 0.5083
Epoch 6/10
9/9 [=====] - 0s 2ms/step - loss: 0.2606 - accuracy: 0.5289
Epoch 7/10
9/9 [=====] - 0s 1ms/step - loss: 0.2603 - accuracy: 0.5248
Epoch 8/10
9/9 [=====] - 0s 2ms/step - loss: 0.2560 - accuracy: 0.6074
Epoch 9/10
9/9 [=====] - 0s 2ms/step - loss: 0.2442 - accuracy: 0.5992
Epoch 10/10
9/9 [=====] - 0s 2ms/step - loss: 0.2316 - accuracy: 0.6281
```

```
Out[41]: <keras.callbacks.History at 0x1d78d06ed70>
```

```
In [42]: model2.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.2314 - accuracy: 0.6066
```

```
Out[42]: [0.23137971758842468, 0.6065573692321777]
```

In [43]: `model2.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 32)	448
dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 8)	136
dense_8 (Dense)	(None, 1)	9

=====
Total params: 1,121
Trainable params: 1,121
Non-trainable params: 0
=====

```
In [44]: model3 = Sequential()
model3.add(Dense(64, input_dim=13, activation='relu'))
model3.add(Dense(32, activation='relu'))
model3.add(Dense(16, activation='relu'))
model3.add(Dense(8, activation='relu'))
model3.add(Dense(1, activation='sigmoid'))
```

```
In [45]: model3.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model3.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 1s 2ms/step - loss: 0.4695 - accuracy: 0.4917
Epoch 2/10
9/9 [=====] - 0s 2ms/step - loss: 0.3186 - accuracy: 0.5702
Epoch 3/10
9/9 [=====] - 0s 2ms/step - loss: 0.2737 - accuracy: 0.6281
Epoch 4/10
9/9 [=====] - 0s 2ms/step - loss: 0.2814 - accuracy: 0.5868
Epoch 5/10
9/9 [=====] - 0s 2ms/step - loss: 0.3501 - accuracy: 0.5289
Epoch 6/10
9/9 [=====] - 0s 2ms/step - loss: 0.2845 - accuracy: 0.5537
Epoch 7/10
9/9 [=====] - 0s 2ms/step - loss: 0.2123 - accuracy: 0.6736
Epoch 8/10
9/9 [=====] - 0s 2ms/step - loss: 0.3161 - accuracy: 0.5992
Epoch 9/10
9/9 [=====] - 0s 2ms/step - loss: 0.2393 - accuracy: 0.6612
Epoch 10/10
9/9 [=====] - 0s 1ms/step - loss: 0.2566 - accuracy: 0.5661
```

```
Out[45]: <keras.callbacks.History at 0x1d78cf7ace0>
```

```
In [46]: model3.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 4ms/step - loss: 0.3637 - accuracy: 0.4918
```

```
Out[46]: [0.36367374658584595, 0.49180328845977783]
```

In [47]: `model3.summary()`

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_9 (Dense)	(None, 64)	896
dense_10 (Dense)	(None, 32)	2080
dense_11 (Dense)	(None, 16)	528
dense_12 (Dense)	(None, 8)	136
dense_13 (Dense)	(None, 1)	9
=====		
Total params: 3,649		
Trainable params: 3,649		
Non-trainable params: 0		

In []: